


Topic 7

Classification: Part 1


One more thing on Topic Modeling

- Is it a type of clustering?
 - Sometimes LDA is referred to as a type of fuzzy clustering
 - But clustering doesn't come up much when discussing topic modeling
 - Soft clustering can be used to try and generate topics, but it isn't quite the same approach as something like LDA, LSA, PLSA, etc.

Classification



Classification



Regression

Pipeline

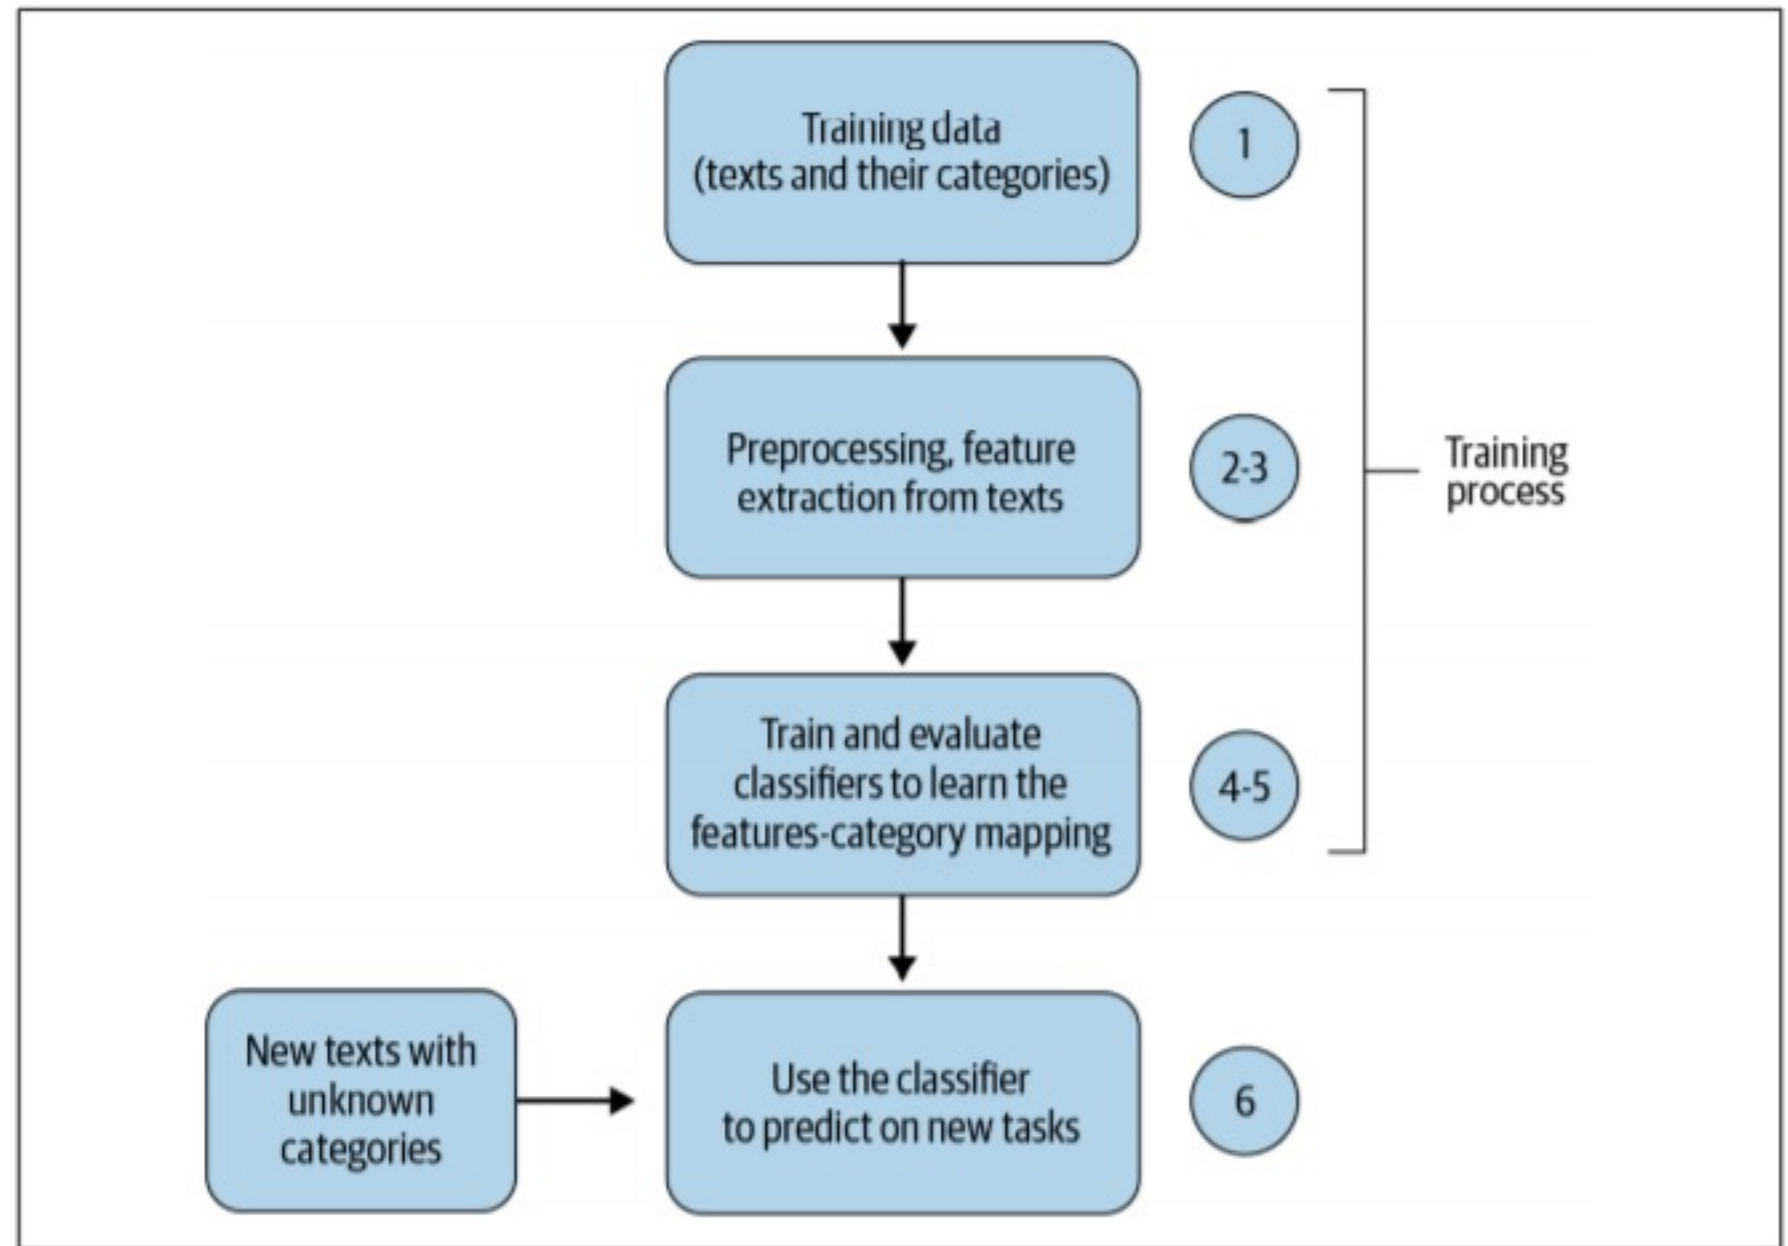




Figure 4-3. Flowchart of a text classification pipeline

Classical NLP

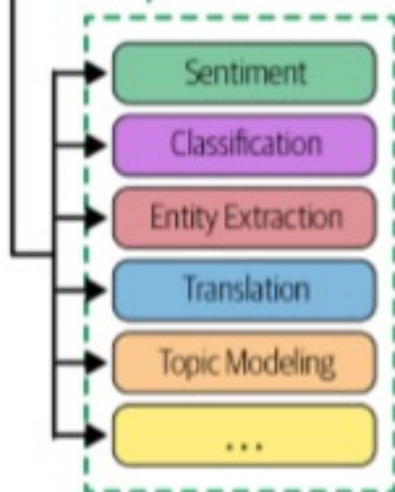
Preprocessing



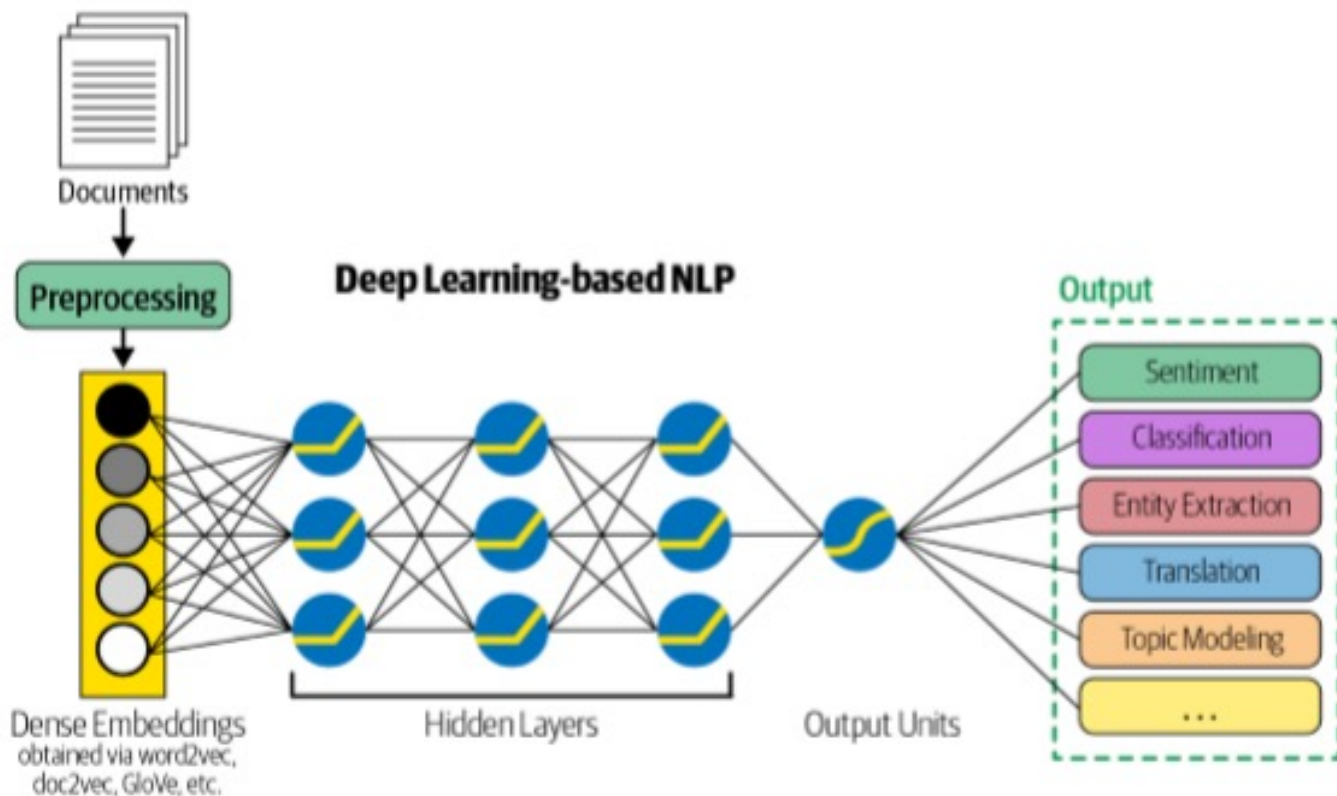
Modeling



Output



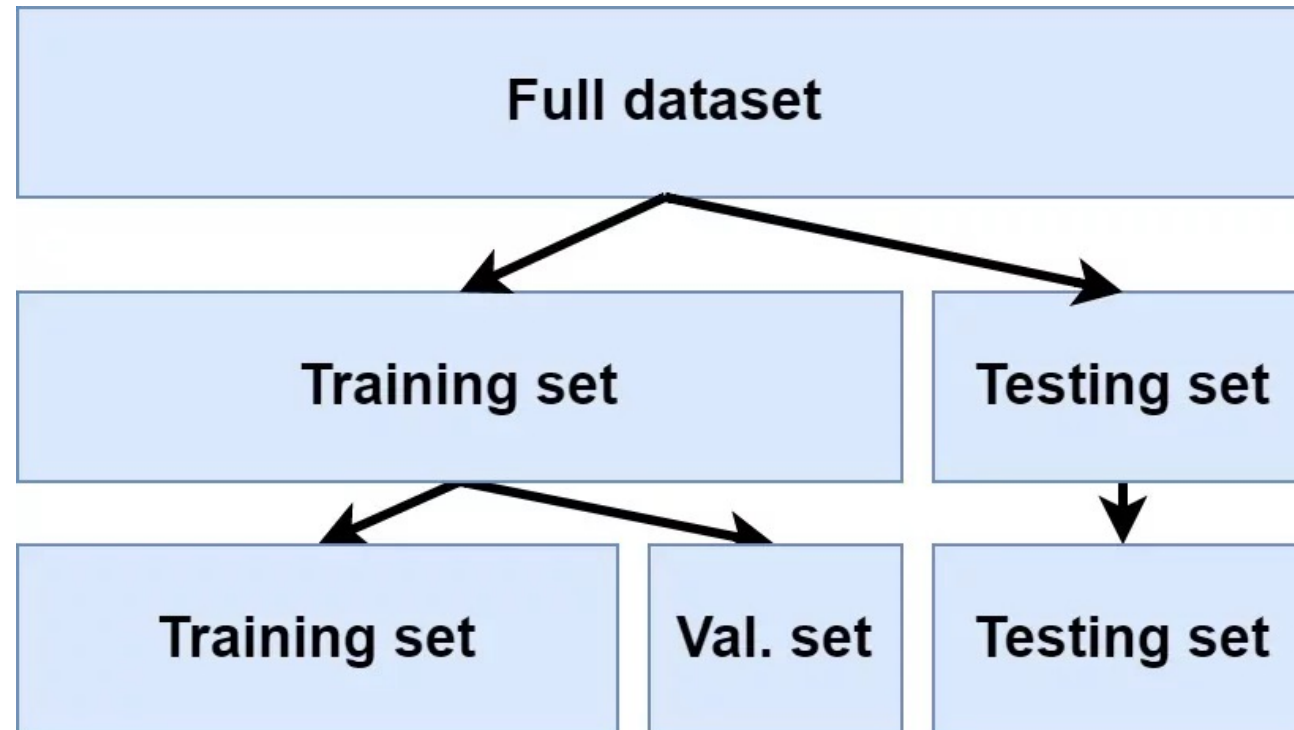
Deep Learning-based NLP



Real-world Examples

- Content classification/organization
- Customer service and Ecommerce
- Author attribution
- Language identification
- Medical research

Data splitting



Generalization error and test error

Generalization error:

- The error made by the model when applied to unseen data

Test set & test error:

- The error made by the model when applied to **a set** of unseen data
- **An unbiased estimate** of generalization error

Training error

- The error made by a model when the model is applied to the data in the training set
- **An over-optimistic estimate** of generalization error

Validation set

- A surrogate for test set during model development

Validation error

- The error made by the model when applied to validation set (implicitly seen by the model)
- A biased estimate of generalization error when there are hyperparameters to tune

Irreducible error

- Irreducible error is the lowest achievable prediction error
- It is a characteristic of the dataset/task under study
- It is independent of the model being used
- It often cannot be calculated analytically
- Often we used human error as an estimate (upper bound) for it

Underfitting

How to identify?

Training error is much higher than the **irreducible error**

How to deal with?

Increasing model complexity



Read more:

[Le, William Trung, et al. "Overview of machine learning: part 2: deep learning for medical image analysis." Neuroimaging Clinics 30.4 \(2020\): 417-431.](#)


Overfitting

How to identify?

Test error is much higher than the **training error**

How to deal with?

- Decrease model complexity
- Collect more data
- Use data augmentation



☠️ **You are there** ☠️

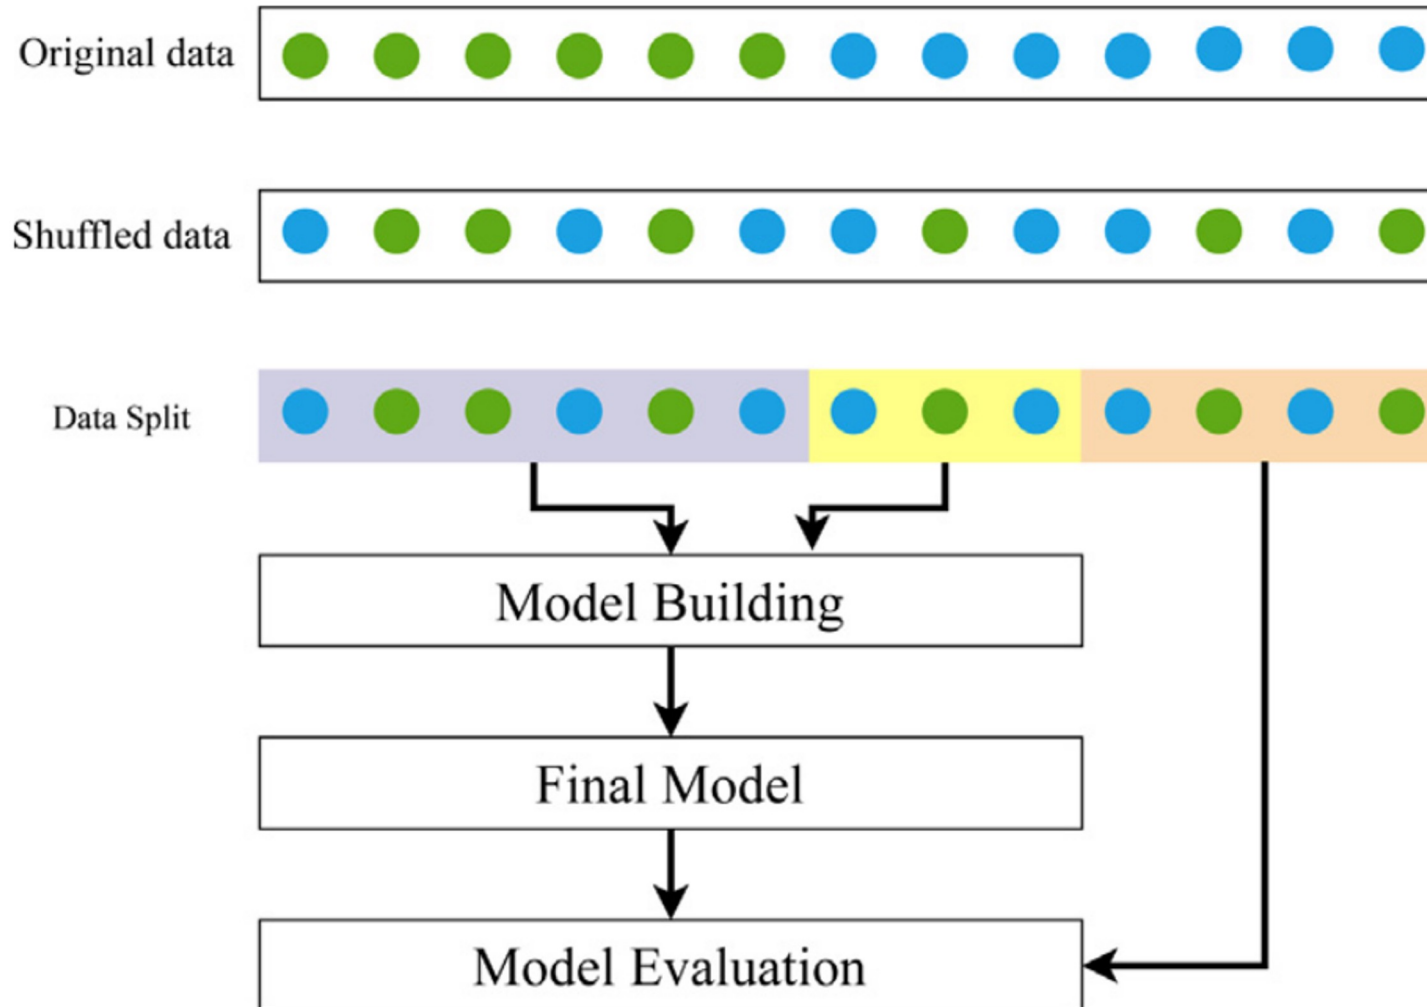
Read more:

[Le, William Trung, et al. "Overview of machine learning: part 2: deep learning for medical image analysis." Neuroimaging Clinics 30.4 \(2020\): 417-431.](#)

Can overfitting happen with Word2Vec?

Calculating training, validation, and test errors

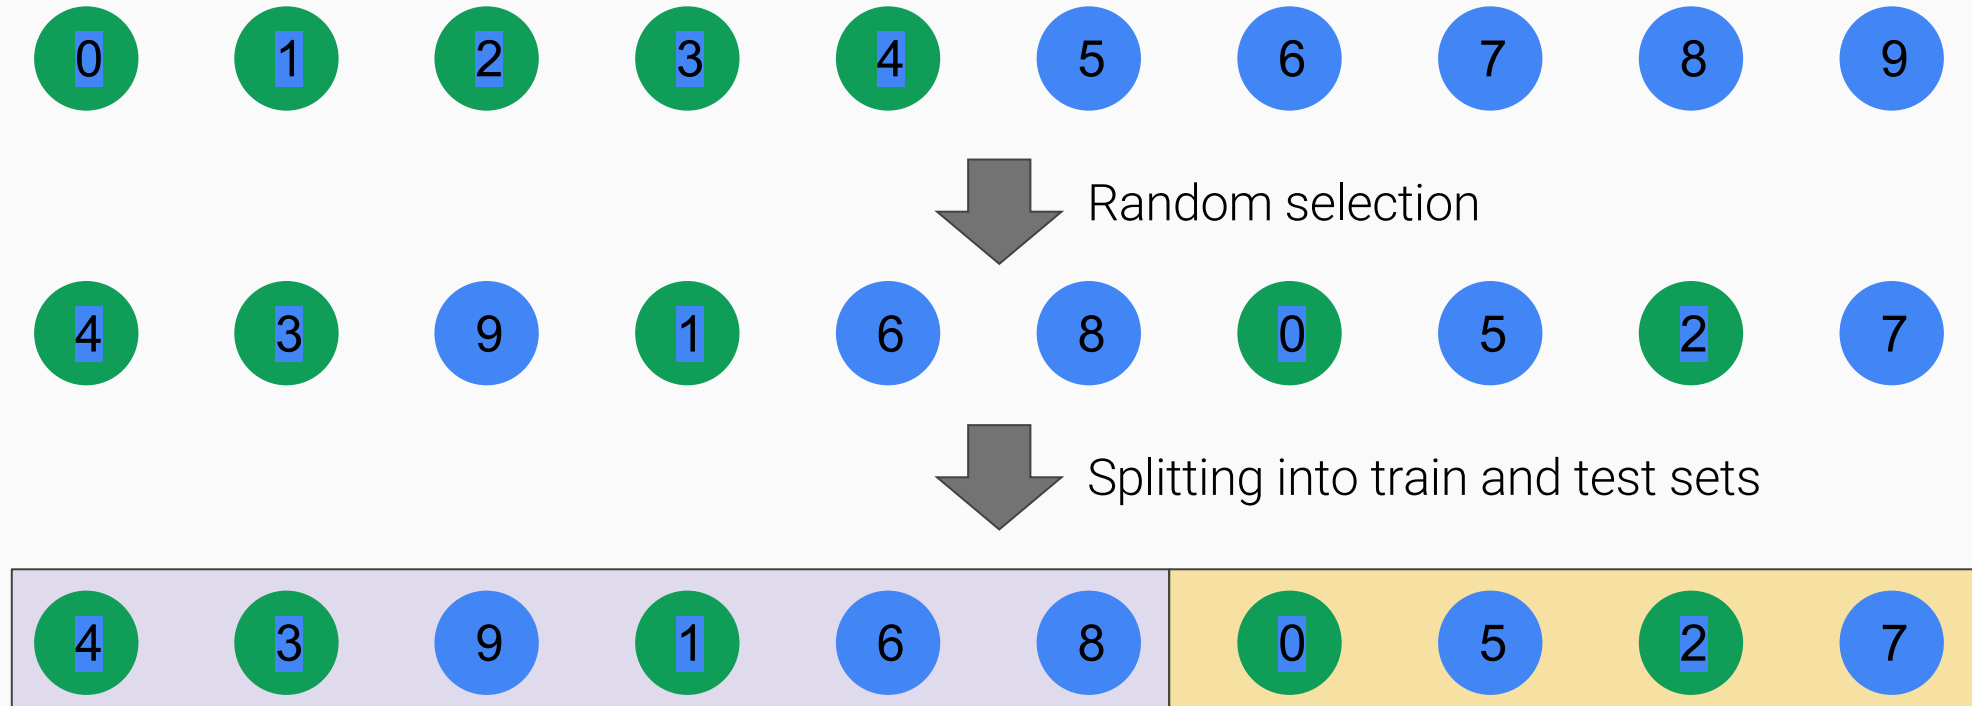
Holdout validation



Why shuffle?

Where is the danger?

Splitting data: Random split



Random split: Code

```
from sklearn.model_selection import train_test_split
train_test_split(*arrays,
                 test_size=None, train_size=None,
                 random_state=None,
                 shuffle=True, stratify=None)
```

This returns sample values.

[Read more](#)

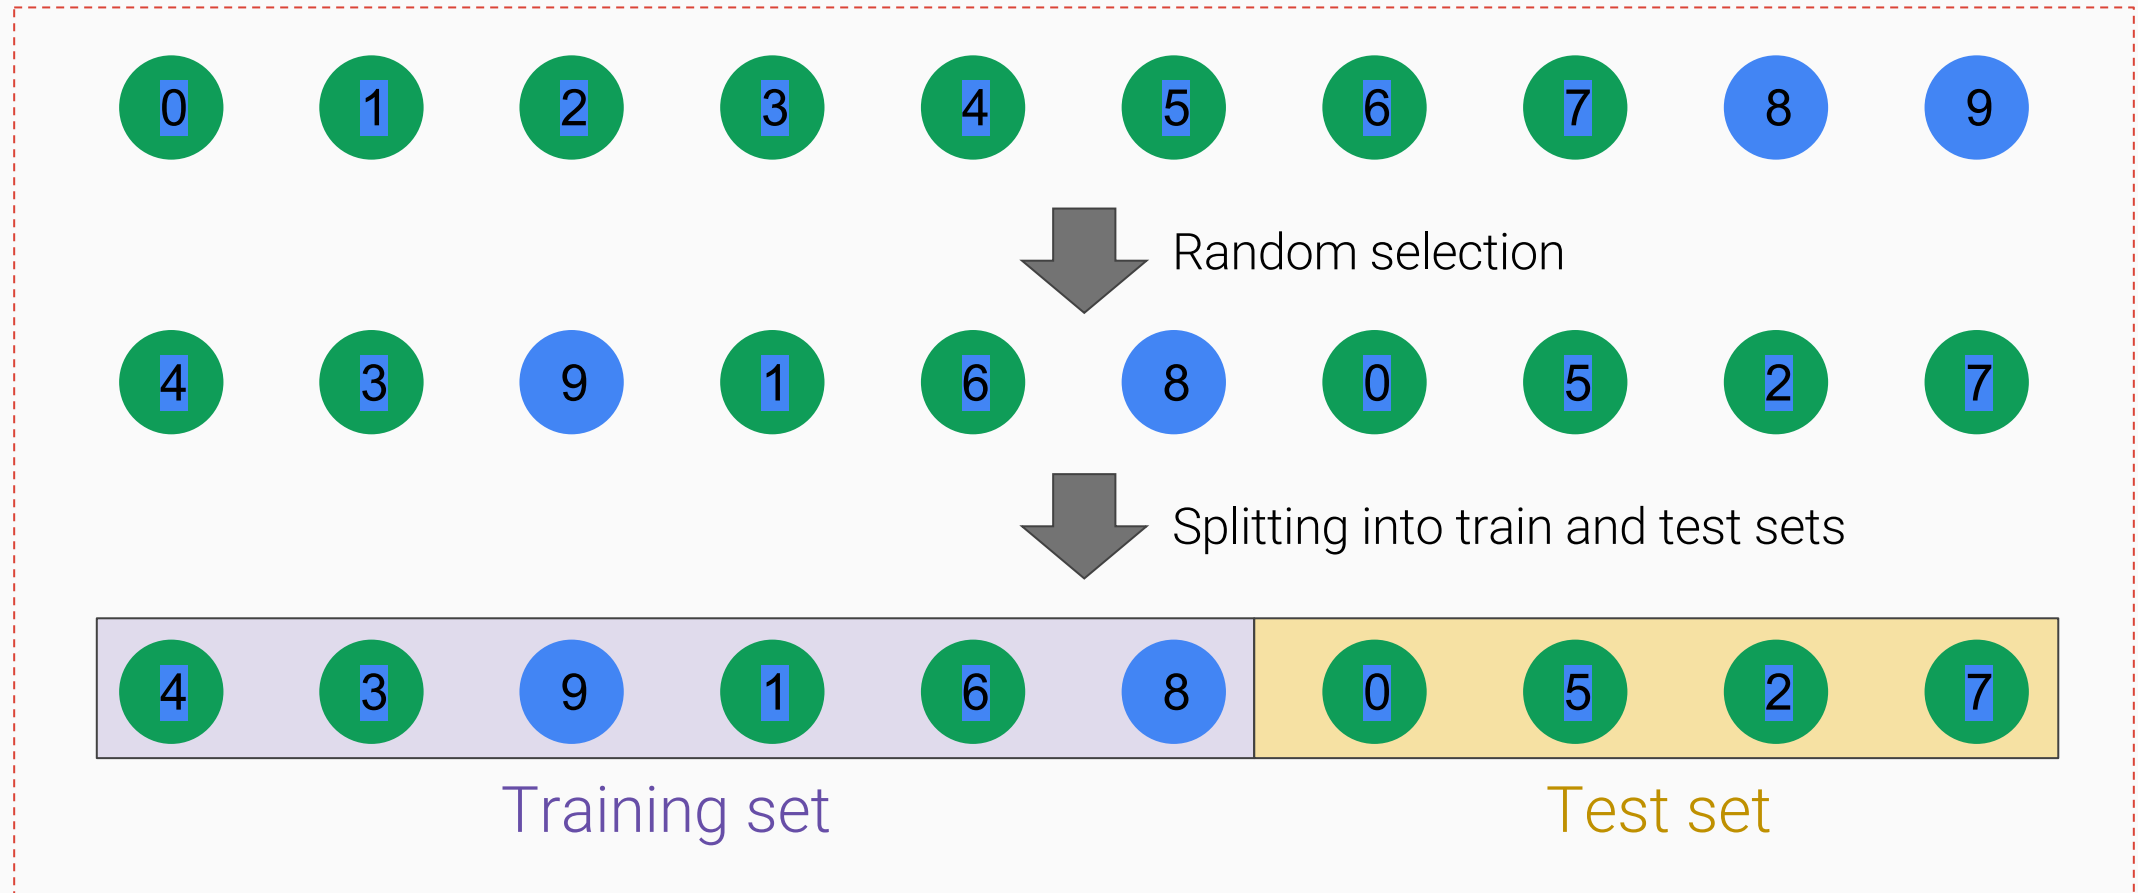
```
from sklearn.model_selection import ShuffleSplit
ShuffleSplit(n_splits=10,
            test_size=None, train_size=None,
            random_state=None)
```

Using the split method, this returns sample indices.

[Read more](#)

What is the main shortcoming of random split?

Class imbalance



Does the test set represent the data distribution?

How to address this issue? Stratified Shuffle Split

```
from sklearn.model_selection import train_test_split
train_test_split(*arrays,
                 test_size=None, train_size=None,
                 random_state=None,
                 shuffle=True,
                 stratify=None)
```

This returns sample values.

[Read more](#)

```
from sklearn.model_selection import StratifiedShuffleSplit
StratifiedShuffleSplit(n_splits=10,
                      test_size=None, train_size=None,
                      random_state=None)
```

Using the split method, this returns sample indices.

[Read more](#)

Group Shuffle Split

Safeguarding the independence of training and test samples

```
from sklearn.model_selection import GroupShuffleSplit
GroupShuffleSplit(n_splits=5,
                  test_size=None, train_size=None,
                  random_state=None)
```

Using the split method, this
returns sample indices.

[Read more](#)

Scenarios

Be aware of the following mistakes:




**Oversampling before
data split**



**Data augmentation
before data split**



**Sample data points
across data splits**



In undersampling, we pull all the rare events while pulling a sample of the abundant events in order to equalize the datasets.

Abundant
dataset

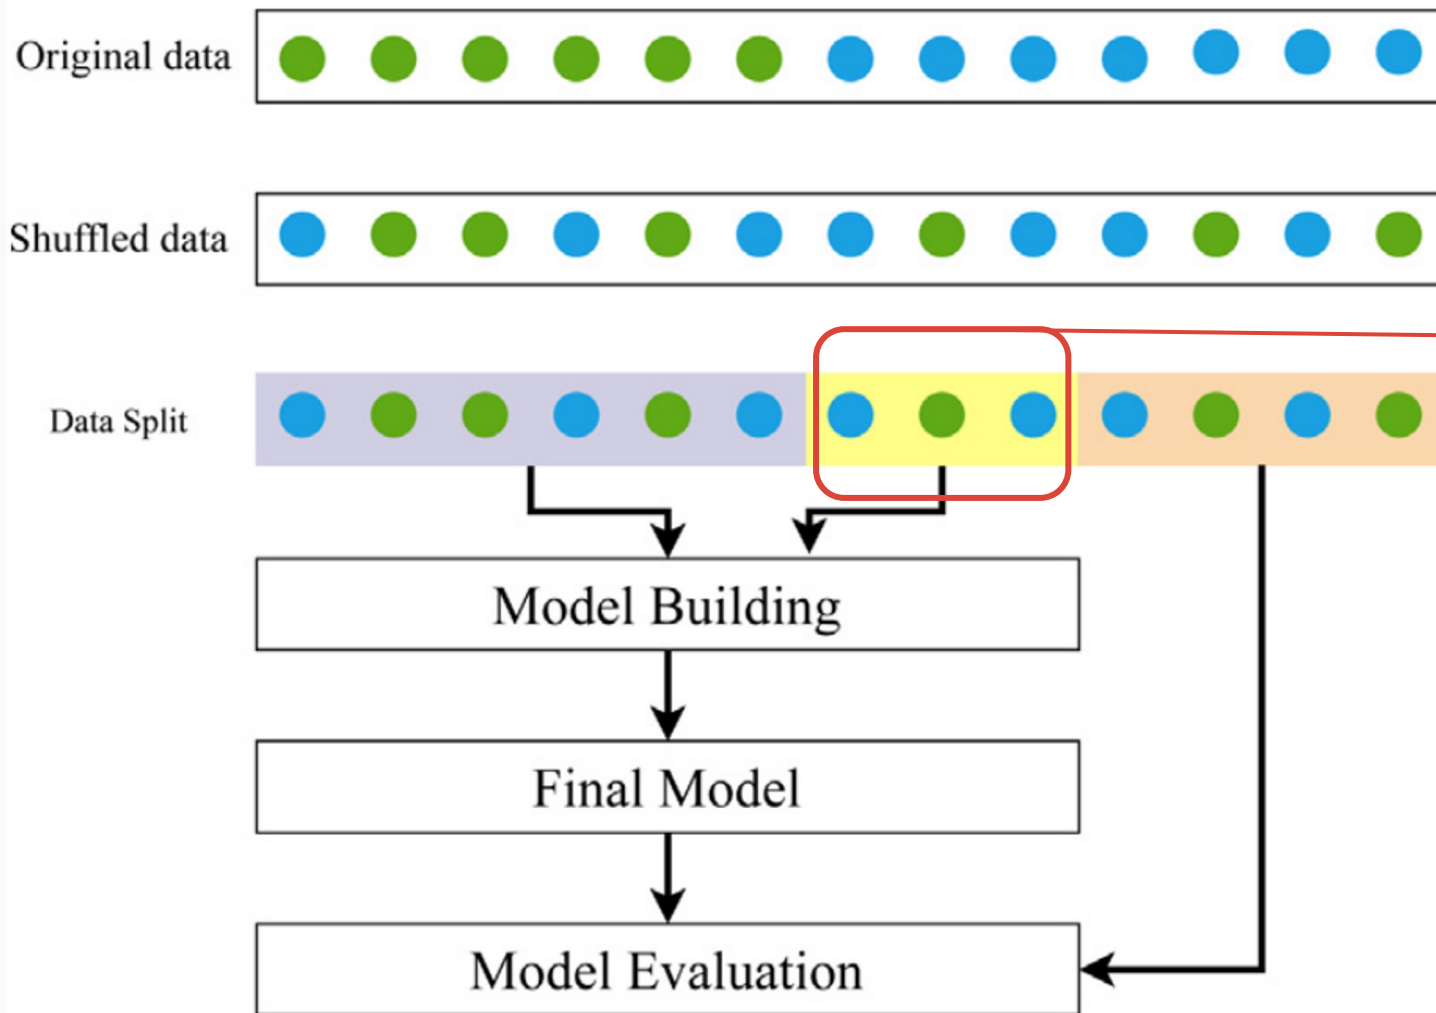
Rare
dataset

These methods can be used separately or together; one is not better than the other. Which method a data scientist uses depends on the dataset and analysis.

Balancing Datasets

Cross-Validation

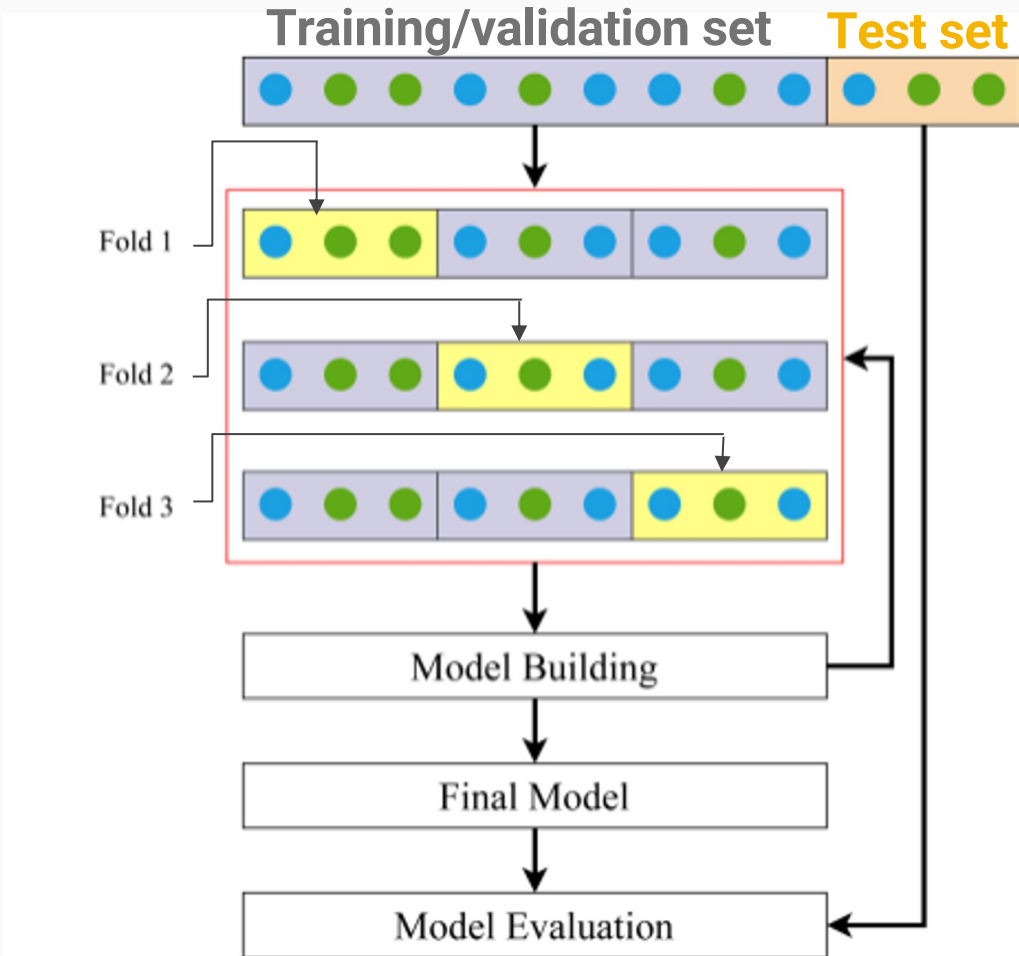
Why cross-validation?



Validation error may significantly vary depending on the composition of **validation set**

Repeating with
different
training and
validation sets

K-fold cross-validation



```
from sklearn.model_selection import KFold
KFold(n_splits=10,
      shuffle=False,
      random_state=None)
```

Using the split method, this returns sample indices. [Read more here!](#)

```
from sklearn.model_selection import StratifiedKFold
StratifiedKFold(n_splits=10,
                shuffle=False,
                random_state=None)
```

Using the split method, this returns sample indices. [Read more here!](#)

```
from sklearn.model_selection import GroupKFold
GroupKFold(n_splits=10,
           random_state=None)
```

Using the split method, this returns sample indices. [Read more here!](#)

Leave-one-out cross-validation

- A special case of K-fold cross-validation

Leave-p-out cross-validation

- An extended form of Leave-one-out cross-validation
- An exhaustive approach and computationally expensive
- Rarely used for $p > 2$

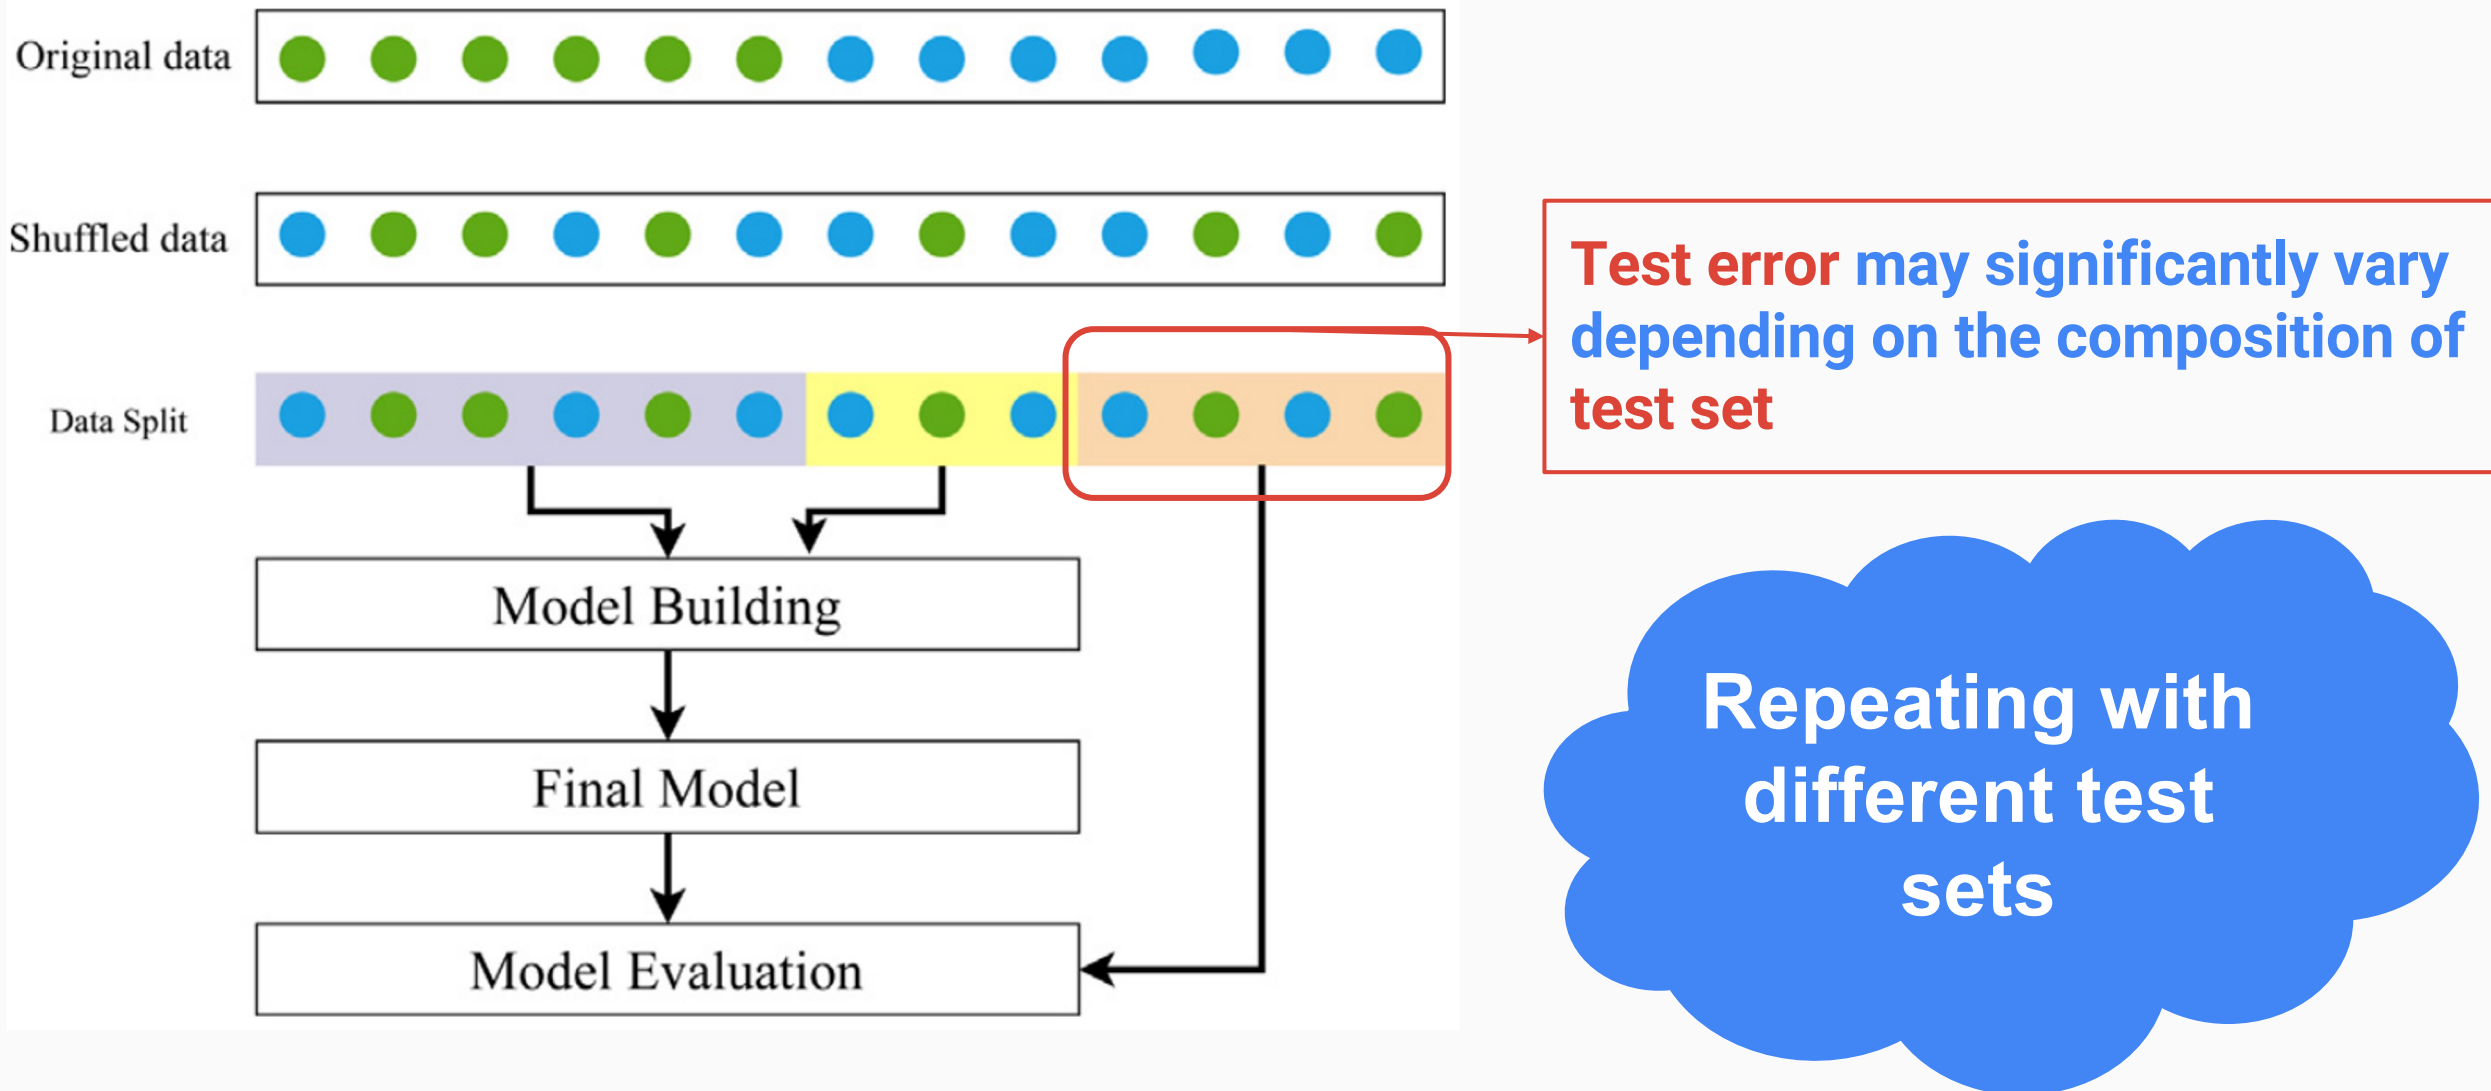
Leave-one-group-out cross-validation

- Similar to one Leave-one-out cross-validation, but uses one group instead of one sample

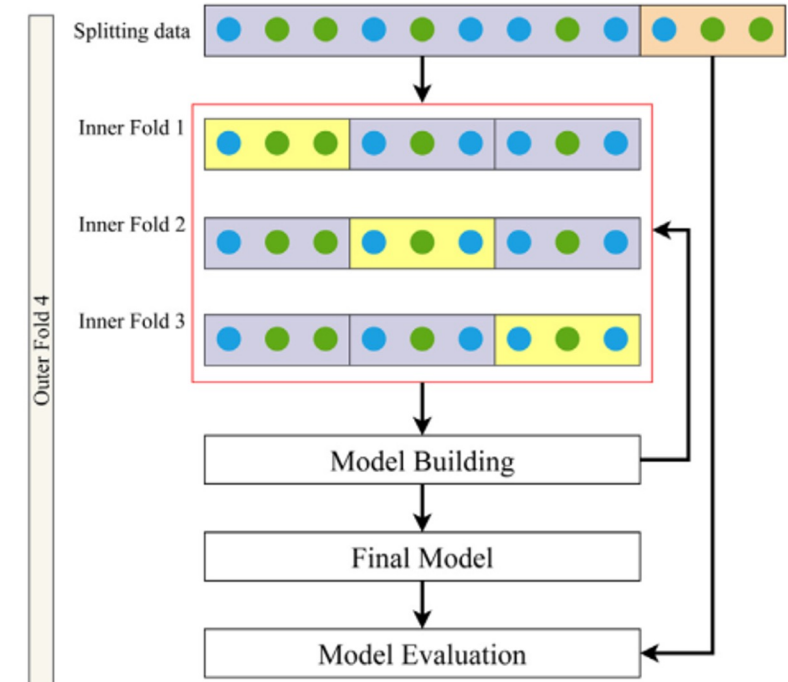
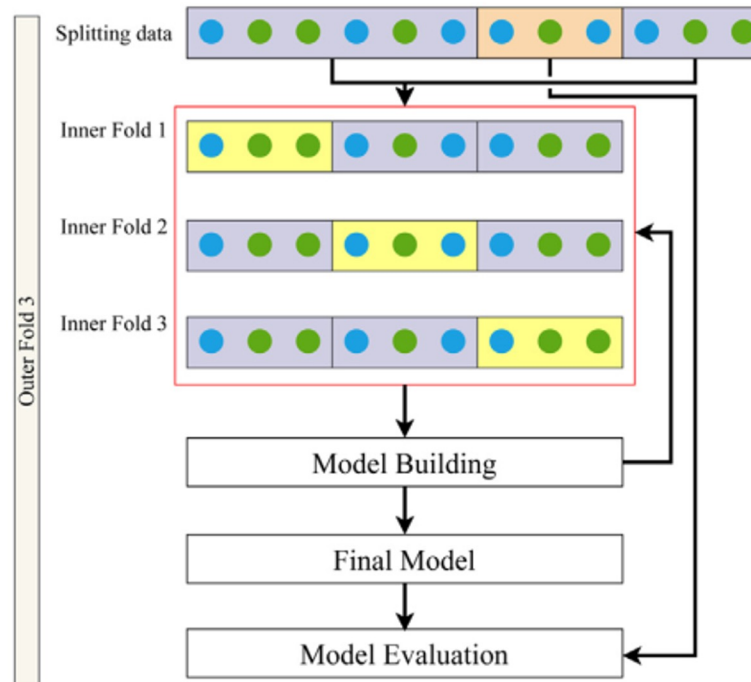
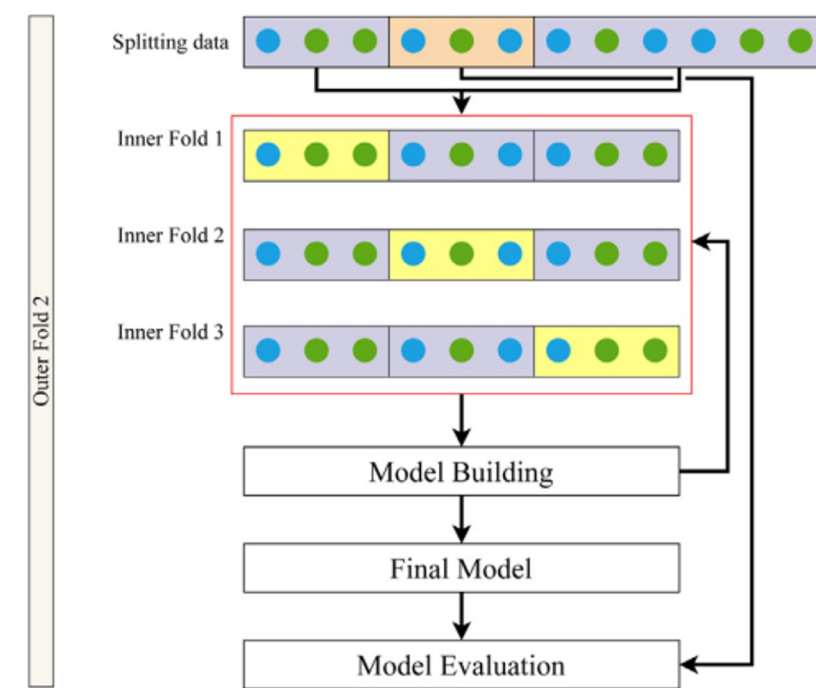
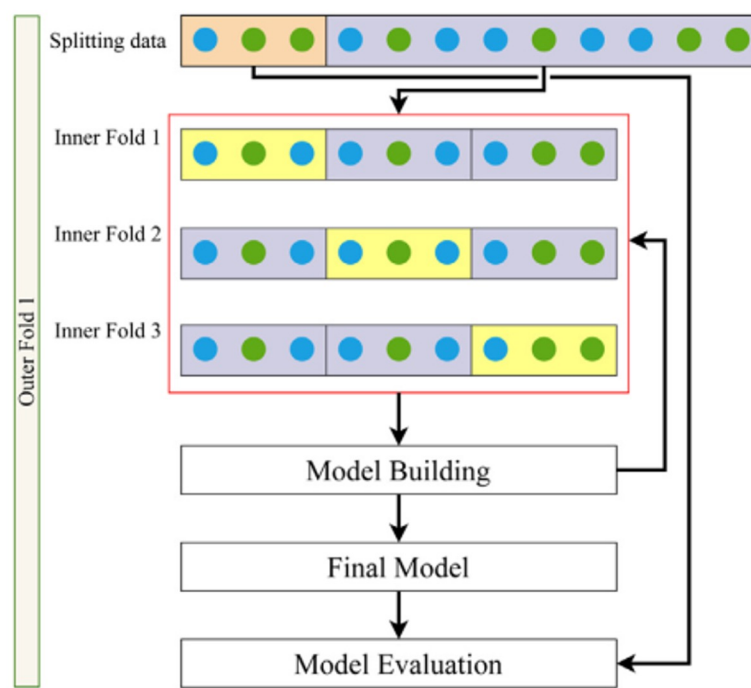
Nested cross-validation

- Applies two cross-validation in a nested manner
- used to train a model in which hyperparameters also need to be optimized


Why nested cross-validation?




A solution



What is a batch effect?



What is a batch effect?



In the context of NLP

- Combining datasets, also referred to as Frankenstein datasets to increase dataset size
 - Maybe the writing style is different depending on platform
 - Age of users or other characteristics

Evaluation Metrics

1. Accuracy Score – no. of correctly classified instances/total no. of instances
2. Precision Score – the ratio of correctly predicted instances over total positive instances
3. Recall Score – the ratio of correctly predicted instances over total instances in that class
4. Roc Curve – a plot of true positive rate against false positive rate
5. Classification Report – report of precision, recall and f1 score
6. Confusion Matrix – a table used to describe the classification models

What does the data look like for classification in NLP?

- Need to perform feature extraction for ML solutions
 - BOW
 - TF-IDF
 - Distributed representation vectors
 - Etc.

Naïve Bayes Classification

- Bayes theorem

$$P(y|X) = \frac{P(X|y) * P(y)}{P(X)}$$

We consider every word in our dataset as independent

$$P(X|y) = P(x_1, x_2, \dots, x_n|y)$$

$$= P(x_1|x_2, \dots, x_n, y) * P(x_2|x_3, \dots, x_n, y) \dots P(x_n|y)$$

$$P(X|y) = P(x_1|y) * P(x_2|y) \dots P(x_n|y)$$

$$P(y|X) = \frac{P(x_1|y) * P(x_2|y) \dots P(x_n|y) * P(y)}{P(x_1) * P(x_2) \dots P(x_n)}$$

Example

15 Not Spam emails and **10 Spam** emails

- $P(\text{Dear} | \text{Not Spam}) = 8/34$
- $P(\text{Visit} | \text{Not Spam}) = 2/34$
- $P(\text{Dear} | \text{Spam}) = 3/47$
- $P(\text{Visit} | \text{Spam}) = 6/47$
- Etc.

	Not Spam	Spam
Dear	8	3
Visit	2	6
Invitation	5	2
Link	2	7
Friend	6	1
Hello	5	4
Discount	0	8
Money	1	7
Click	2	9
Dinner	3	0
Total Words	34	47

$$P(\text{Hello Friend}|\text{Not Spam}) = P(\text{Hello}|\text{Not Spam}) * P(\text{Friend}|\text{Not Spam})$$

$$P(\text{Not Spam}|\text{Hello Friend}) = P(\text{Hello}|\text{Not Spam}) * P(\text{Friend}|\text{Not Spam}) * P(\text{Not Spam})$$

$$P(\text{Not Spam}|\text{Hello Friend}) = \frac{5}{34} * \frac{6}{34} * \frac{15}{25} = 0.0155$$

$$P(\text{Spam}|\text{Hello Friend}) = \frac{4}{47} * \frac{1}{47} * \frac{10}{25} = 0.00072$$

Laplacian Smoothing

$$\begin{aligned} &P(\text{Not Spam} | \text{dear visit dinner money money money}) \\ &= P(\text{dear visit dinner money money money} | \text{Not Spam}) * P(\text{Not Spam}) \end{aligned}$$

$$\begin{aligned} &P(\text{Spam} | \text{dear visit dinner money money money}) \\ &= P(\text{dear visit dinner money money money} | \text{Spam}) * P(\text{Spam}) \end{aligned}$$

$$\hat{\theta} = \frac{x_i + \alpha}{N + \alpha d} \quad (i = 1, \dots, d)$$

How can we improve?

- Maybe we are using too many features
- Few samples of relevant articles (class imbalance)
- Better algorithm
- Preprocessing and feature extraction
- Hyperparameter tuning

Many other options for classification

- SVM
- Logistic Regression
- KNN
- Decision tree
- Random Forest
- XGBoost
- etc.

Next time

- Pre-trained models
- Using these pre-trained models for classification